
2. ParallelDemo

February 16, 2014

Part I

Sample Codes

NOTE:

Each process has an associated identifier. The process providing the interactive julia prompt always has an id equal to 1, as would the julia process running the driver script in the example above. The processors used by default for parallel operations are referred to as workers. When there is only one process, process 1 is considered a worker. Otherwise, workers are considered to be all processes other than process 1.

The base Julia installation has in-built support for two types of clusters:

1. A local cluster specified with the `-p` option as shown above. 2. And a cluster spanning machines using the `-machinefile` option. This uses ssh to start the worker processes on the specified machines.

Functions `addprocs`, `rmprocs`, `workers` and others, are available as a programmatic means of adding, removing and querying the processes in a cluster.

```
workers()
In [1]: 1-element Array{Int64,1}:
Out [1]: 1
```

```
procs()
In [2]: 1-element Array{Int64,1}:
Out [2]: 1
```

```
addprocs(8)
In [3]: 8-element Array{Any,1}:
Out [3]: 2
         3
         4
         5
         6
         7
         8
         9
```

```
r = remotecall(3, rand, 2, 2)
In [4]: RemoteRef{3,1,4}
Out [4]:
```

```

b=fetch(r)
In [5]: 2x2 Array{Float64,2}:
Out [5]:  0.525251  0.775677
          0.440102  0.352892

s = @spawnat 2 rand(2,2)
In [6]: RemoteRef(2,1,6)
Out [6]: fetch(s)
In [7]: 2x2 Array{Float64,2}:
Out [7]:  0.243035  0.296231
          0.405337  0.618076

remotecall_fetch(2, rand, 2,2)
In [8]: 2x2 Array{Float64,2}:
Out [8]:  0.668564  0.0206293
          0.0210516  0.774269

remotecall_fetch(2, getindex,r,1,1)
In [9]: 0.5252510559469807
Out [9]: r = @spawn rand(2,2)
In [10]: RemoteRef(3,1,10)
Out [10]: fetch(r)
In [11]: 2x2 Array{Float64,2}:
Out [11]:  0.164833  0.311987
          0.604117  0.707194

s = @spawn 1+fetch(r)
In [12]: RemoteRef(3,1,12)
Out [12]: fetch(s)
In [13]: 2x2 Array{Float64,2}:
Out [13]:  1.16483  1.31199
          1.60412  1.70719

a=RemoteRef[]
In [14]: for i in 1:10
          push!(a,@spawn sin(i))
        end

@parallel (+) for i in a
In [15]: fetch(i)
        end
1.4111883712180104

Out [15]: sum(map(x->fetch(x),a))
In [16]: 1.4111883712180104
Out [16]: @everywhere function producer()
In [21]:   for i in 1:3
          produce(i)
          end
          end
          task=Task(producer)
          consume(task)
1

Out [21]: consume(task)
In [18]: 2
Out [18]:

```

```

In [24]: n=10;
          @everywhere function ssin()
            for i in 1:n
              produce(sin(i))
            end
          end
          p=Task(ssin)
          reduce(+,map(x->consume(p),1:n))
1.4111883712180104

Out [24]: s=sin(1)
In [26]: @time for i in 2:10000000
          s+=sin(i)
        end
        s
elapsed time: 2.701724951 seconds (319999968 bytes allocated)
1.9558914085412367

Out [26]: @time @parallel (+) for i=1:10000000
In [27]:   sin(i)
        end
elapsed time: 0.476710806 seconds (517064 bytes allocated)
1.9558914085412054

Out [27]: @time reduce(+,(map(i->sin(i),1:10000000)))
In [28]: elapsed time: 3.724218083 seconds (879985408 bytes allocated)
1.9558914085412236

Out [28]: @time reduce(+,[sin(i) for i in 1:10000000])
In [29]: elapsed time: 2.61912771 seconds (400000032 bytes allocated)
1.9558914085412236

Out [29]: @time sum( map(1:10000000) do x
In [30]:   sin(x);
        end
        )
elapsed time: 3.178096776 seconds (559984848 bytes allocated)
1.9558914085412236

Out [30]:

```

Part II

Data Movement

method 1 $A = \text{rand}(1000,1000)$ $Bref = @\text{spawn } A^2 \dots \text{fetch}(Bref)$

method 2 $Bref = @\text{spawn } \text{rand}(1000,1000)^2 \dots \text{fetch}(Bref)$

Part III

Parallel Maps/Loops

```

In [47]: @everywhere function count_heads(n)
          c::Int = 0
          for i=1:n
            c += randbool()
          end
          c

```

```
end
@time count_heads(200000000)/200000000
```

In [48]: elapsed time: 0.688498406 seconds (159000 bytes allocated)
0.50000936

Out [48]: `#require("count_heads")`

```
In [49]: @time begin
a = @spawn count_heads(100000000)
b = @spawn count_heads(100000000)
(fetch(a)+fetch(b))/200000000
end
elapsed time: 0.395470371 seconds (21752 bytes allocated)
0.50006082
```

Out [49]: `@time begin`

```
In [50]: nheads = @parallel (+) for i=1:200000000
int(randbool())
end
nheads/200000000
end
elapsed time: 0.168353892 seconds (526264 bytes allocated)
0.49999534
```

Out [50]: `@time begin`

```
In [51]: a=0
for i=1:200000000
a+=int(randbool())
end
a/200000000
end
elapsed time: 8.610191092 seconds (3199988192 bytes allocated)
0.50003185
```

Out [51]: `@time sum(map(x->int(randbool()),1:200000000))/200000000`

In [54]: elapsed time: 22.731768799 seconds (7999985040 bytes allocated)
0.49999467

Out [54]: `b = zeros(100000);`

```
In [55]: @parallel for i=1:100000
b[i] = i
end
```

```
b[1:10] // assignment is local to the process in @parallel for
```

In [56]: 10-element Array{Float64,1}:

Out [56]: 0.0
0.0
0.0
0.0
0.0
0.0
0.0
0.0
0.0
0.0