
3. Distributed Arrays and Map-Reduce

February 13, 2014

```
addprocs(8)

In [1]: 8-element Array{Any,1}:
Out [1]: 2
         3
         4
         5
         6
         7
         8
         9

nprocs()

In [2]: 9
Out [2]: dat=map(x->integer(x*10),rand(10000,10000))
In [3]: 10000x10000 Array{Int64,2}:
Out [3]:  7  9  3  6  0  3  3  3  0  9  6  ...  0  7  7  9 10  2  1
          3  3
          8 10  6  1  5  2  5  7  4  9  3  ...  6  7  2 10  5  6  2
          2  9
          6  5  8  8  3  6  1  9  5  6  4  ...  4  7  9  6  1  1  4
          9  5
          2  3  5  8 10  1  7  7  9  7  5  ...  4  7  1  1 10  9  4
          6  7
          0  0  1  3  2  8  9 10  6  7  1  ...  5  2  3  3  1  8  6
          1  8
          7  7  6  2  5  9  9  9  7  3  0  ...  3 10  3  2 10  7  4
          3  7
          4  3  4  5  7  8  8  9  8  0  5  ... 10  6  9  4  6  7  9
          2  7
          1  2 10  0  4  6  1  2  3  1  7  ...  1  0  1  3  7  5  2
          4  8
          9  4  4  0  8  3  8  7  9  1  5  ...  7  6  8  5  1  7  5
          2  3
          8  9  8  8  0  8  8  9  4  0  3  ...  1  1  2  8  9  3  0
          9  8
          5  0 10  9  9  6  8  8 10  7  6  ...  7 10  6  7  8  4  1
          4  4
          8  6 10  0  6  4  5  7  9  6  0  ...  9  0  4  4  4  3  9
          5  9
          0  6  9  3  2  9  6  7  1  9  3  ... 10  2  1  7  5  2  5
          6  4
          ⋮          ⋮          ⋮  ...          ⋮
```

```

6 8 6 8 2 6 9 7 3 3 5 5 2 2 1 6 2 8
6 2
6 7 5 1 6 7 3 4 2 6 7 8 3 7 4 1 2 8
3 1
7 3 3 5 7 9 7 8 0 1 7 ... 3 7 7 2 0 4 3
5 9
7 2 3 0 5 6 2 3 1 9 6 0 4 8 9 4 9 6
5 10
5 8 2 7 1 1 2 5 5 2 1 6 4 3 8 1 1 6
9 5
8 1 5 3 3 3 0 6 6 9 8 1 5 7 2 6 6 8
7 6
0 4 0 1 2 7 7 4 9 7 2 5 2 8 1 9 8 7
10 9
6 3 2 2 8 4 8 8 8 4 8 ... 7 4 10 5 6 5 8
2 8
1 3 0 1 1 9 0 6 10 4 8 8 6 5 2 0 4 1
9 2
5 1 4 8 9 4 1 1 3 7 10 5 10 7 3 7 6 3
2 1
6 7 8 7 6 4 4 7 5 6 4 0 4 9 3 6 8 5
7 9
8 2 8 6 10 9 6 10 4 6 8 2 4 1 4 3 6 2
1 4

```

```
ddat=distribute(dat)
```

```

In [4]: 10000x10000 DArray{Int64,2,Array{Int64,2}}:
Out [4]: 7 9 3 6 0 3 3 3 0 9 6 ... 0 7 7 9 10 2 1
3 3
8 10 6 1 5 2 5 7 4 9 3 6 7 2 10 5 6 2
2 9
6 5 8 8 3 6 1 9 5 6 4 4 7 9 6 1 1 4
9 5
2 3 5 8 10 1 7 7 9 7 5 4 7 1 1 10 9 4
6 7
0 0 1 3 2 8 9 10 6 7 1 5 2 3 3 1 8 6
1 8
7 7 6 2 5 9 9 9 7 3 0 ... 3 10 3 2 10 7 4
3 7
4 3 4 5 7 8 8 9 8 0 5 10 6 9 4 6 7 9
2 7
1 2 10 0 4 6 1 2 3 1 7 1 0 1 3 7 5 2
4 8
9 4 4 0 8 3 8 7 9 1 5 7 6 8 5 1 7 5
2 3
8 9 8 8 0 8 8 9 4 0 3 1 1 2 8 9 3 0
9 8
5 0 10 9 9 6 8 8 10 7 6 ... 7 10 6 7 8 4 1
4 4
8 6 10 0 6 4 5 7 9 6 0 9 0 4 4 4 3 9
5 9
0 6 9 3 2 9 6 7 1 9 3 10 2 1 7 5 2 5
6 4
:
:
: ...
:

```

```
map(x->myid(), dat)
```

```
Out [5]: 1 1 1 1 1 1 1 1 1 1 1 1 1 ... 1 1 1 1 1 1 1 1 1
```

```
map(x->myid(), ddat)
```

```
Out [6]: 2 2 2 2 2 2 2 2 2 2 2 2 2 ... 8 8 8 8 8 8 8 8 8
```

```

3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 ... 9 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 ... 9 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9
3 3 3 3 3 3 3 3 3 3 3 3 3 3 9 9 9 9 9 9 9 9
9 9 9

```

```
@time map(sin,dat);
```

In [6]: elapsed time: 18.495849369 seconds (4000237720 bytes allocated)

```
@time map(sin,ddat);
```

In [6]: elapsed time: 3.801187871 seconds (1143612 bytes allocated)

```

In [7]: function intsum(x)
        if x==0
            0
        else
            x+intsum(x-1)
        end
    end
    intsum (generic function with 1 method)

```

Out [7]: @time reduce(+,map(intsum,dat))

In [8]: elapsed time: 13.26386314 seconds (3573678796 bytes allocated)
1925337529

Out [8]: @everywhere function intsum(x)

```

In [9]: if x==0
        0
    else
        x+intsum(x-1)
    end
end

```

```
@time reduce(+,map(intsum,ddat))
```

In [10]: elapsed time: 6.24797383 seconds (1743872 bytes allocated)
1925337529

Out [10]: @everywhere function rsumtail(x,n)

```

In [11]: if x==0
        n
    else

```

```

    rsumtail(x-1,n+x)
  end
end

@everywhere function rsumtail(x)
  rsumtail(x,0)
end

@everywhere function rsum(x)
  if x==0
    0
  else
    x+rsum(x-1)
  end
end

@time rsumtail(100000)

```

In [12]: elapsed time: 0.00278938 seconds (42804 bytes allocated)
5000050000

Out [12]: @time rsum(100000)

In [13]: elapsed time: 0.002768959 seconds (23496 bytes allocated)
5000050000

Out [13]: @time (s=0;

In [14]: for i in dat
s+=intsum(i)
end;
s)

elapsed time: 26.858735368 seconds (6400048908 bytes allocated)
1925337529

Out [14]: @time @parallel (+) for i in dat

In [21]: intsum(i)
end

elapsed time: 30.597089961 seconds (8590437940 bytes allocated)
1925337529

Out [21]: @time reduce(+,map(intsum,ddat))

In [16]: elapsed time: 7.447626368 seconds (1453340 bytes allocated)
1925337529

Out [16]: ddat.indexes

In [22]: 2x4 Array{Range{Int64},2}:

Out [22]: (1:5000,1:2500) (1:5000,2501:5000) ... (1:5000,7501:10000)
(5001:10000,1:2500) (5001:10000,2501:5000)
(5001:10000,7501:10000)

ddat.chunks

In [23]: 2x4 Array{RemoteRef,2}:

Out [23]: RemoteRef(2,1,5) RemoteRef(4,1,7) RemoteRef(6,1,9)
RemoteRef(8,1,11)
RemoteRef(3,1,6) RemoteRef(5,1,8) RemoteRef(7,1,10)
RemoteRef(9,1,12)

ddat.pmap

In [24]: 8-element Array{Int64,1}:

Out [24]: 2
3
4
5
6

```

7
8
9
fetch(@spawnat 2 sum(x->x*x,localpart(ddat)))
In [27]: 418800280
Out [27]: map(fetch,{@spawnat p sum(map(x->x*x,localpart(ddat))) for p in ddat.pmap})
In [28]: 8-element Array{Any,1}:
Out [28]: 418800280
          418938675
          418841631
          418547283
          419008882
          418728812
          418788879
          418975117

@time reduce(+,map(fetch,{@spawnat p sum(map(x->x*x,localpart(ddat))) for p in procs(d
In [29]: elapsed time: 2.15132193 seconds (663104 bytes allocated)
          3350629559
Out [29]: @time reduce(+,map(x->x*x,ddat))
In [39]: elapsed time: 3.018259173 seconds (1472940 bytes allocated)
          3350629559
Out [39]: @time sum(map(x->x*x,dat))
In [42]: elapsed time: 8.165484936 seconds (2399993056 bytes allocated)
          3350629559
Out [42]:

```