
7. Machine Learning

February 16, 2014

```
In [112]: import RDatasets
import DataFrames
import Clustering
import SVM
import DecisionTree
```

```
In [112]: # read data
iris = RDatasets.data("datasets", "iris");
features = matrix(iris[:, 2:5]);
labels = matrix(iris[:, 6]);
sz, ft = size(features);
train = randbool(sz)
test = !train;
```

1 K-means Clustering

```
In [112]: nc=3
km=Clustering.kmeans(features', nc);
centers=km.centers;
```

Iters	objv	objv-change	affected
1	9.215602e+01	-5.825398e+01	2
2	8.368890e+01	-8.467122e+00	3
3	7.986398e+01	-3.824915e+00	2
4	7.919714e+01	-6.668418e-01	2
5	7.885144e+01	-3.457012e-01	0
6	7.885144e+01	0.000000e+00	0

K-means converged with 6 iterations (objv = 78.85144142614597)

```
km.centers
```

```
In [113]: 4x3 Array{Float64,2}:
```

```
Out [113]: 5.90161  5.006  6.85
2.74839  3.428  3.07368
4.39355  1.462  5.74211
1.43387  0.246  2.07105
```

2 SVM

```
In [114]: Y = [species == "setosa" ? 1.0 : -1.0 for species in iris[:, "Species"]];
X=features'
model=SVM.svm(X[:,train],Y[train])
prediction=SVM.predict(model, X[:,test])
DecisionTree.confusion_matrix(prediction,Y[test])
```

```

Classes:  [-1.0,1.0]
Out [114]: Matrix:  2x2 Array{Int64,2}:
           50   0
           0  22
Accuracy:  1.0
Kappa:    1.0

```

3 Random Forest

```

In [114]: modelrf=DecisionTree.build_forest(labels[train],features[train,:],4,30);
          predrf=DecisionTree.apply_forest(modelrf,features[test,:]);
          println(DecisionTree.confusion_matrix(labels[test],predrf));

```

```

Classes:  {"setosa","versicolor","virginica"}
Matrix:   3x3 Array{Int64,2}:
           22   0   0
           1  24   2
           0   3  20
Accuracy:  0.9166666666666666
Kappa:    0.8745280278826603

```

```

In [114]: accuracyrf = DecisionTree.nfoldCV_forest(labels, features, 4, 30, 3);
          Fold 1

```

```

Classes:  {"setosa","versicolor","virginica"}
Matrix:   3x3 Array{Int64,2}:
           17   0   0
           0  15   0
           0   1  17
Accuracy:  0.98
Kappa:    0.96996996996997

```

```

Fold 2
Classes:  {"setosa","versicolor","virginica"}
Matrix:   3x3 Array{Int64,2}:
           16   0   0
           0  16   1
           0   3  14
Accuracy:  0.92
Kappa:    0.879951980792317

```

```

Fold 3
Classes:  {"setosa","versicolor","virginica"}
Matrix:   3x3 Array{Int64,2}:
           17   0   0
           0  16   2
           0   1  14
Accuracy:  0.94
Kappa:    0.9099099099099098

```

```

Mean Accuracy: 0.9466666666666667

```

4 Adaptive-Boosted Tree (Adaboost)

```

modelad, coeffs=DecisionTree.build_adaboost_stumps(labels[train], features[train,:], 50);
In [114]: predad=DecisionTree.apply_forest(modelad, features[test,:]);
println(DecisionTree.confusion_matrix(labels[test], predad));
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 22  0  0
  1  8 18
  0  0 23
Accuracy: 0.7361111111111112
Kappa: 0.6112531969309464

accuracyad = DecisionTree.nfoldCV_stumps(labels, features, 50, 3);
In [114]: Fold 1
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 0 22  0
 0 18  0
 0  0 10
Accuracy: 0.56
Kappa: 0.34523809523809534

Fold 2
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
13  0  0
 0  5 13
 0  0 19
Accuracy: 0.74
Kappa: 0.6019595835884874

Fold 3
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
15  0  0
 0 10  4
 0  1 20
Accuracy: 0.9
Kappa: 0.8453927025355595

Mean Accuracy: 0.7333333333333334

```

5 Pruned Tree

```

modelpt=DecisionTree.build_tree(labels[train], features[train,:])
In [114]: predpt=DecisionTree.apply_tree(modelpt, features[test,:]);
println(DecisionTree.confusion_matrix(labels[test], predpt));
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 22  0  0
  1 25  1
  0  3 20
Accuracy: 0.9305555555555556
Kappa: 0.8953184065135215

```

```

accuracyp = DecisionTree.nfoldCV_tree(labels, features, 0.9, 3);
In [114]: Fold 1
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 20  0  0
  0 14  1
  0  0 15
Accuracy: 0.98
Kappa: 0.9696969696969696

Fold 2
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 15  0  0
  0 15  0
  0  1 19
Accuracy: 0.98
Kappa: 0.9697885196374622

Fold 3
Classes: {"setosa", "versicolor", "virginica"}
Matrix: 3x3 Array{Int64,2}:
 15  0  0
  3 15  2
  0  1 14
Accuracy: 0.88
Kappa: 0.8203592814371259

Mean Accuracy: 0.9466666666666667

```

6 Neural Network

```

In [8]: # https://github.com/nwenzel/Julia\_Neural\_Network
# Code by Nathan Wenzel

function sigmoid(z)
    # sigmoid is a basic sigmoid function returning values from 0-1
    1. / ( 1. + 1e.^(-z) )
end

function sigmoidGradient(z)
    sigmoid(z) .* ( 1 - sigmoid(z) )
end

function initialize_theta(input_unit_count, output_class_count, hidden_unit_length_list)
    #
    # initialize_theta creates architecture of neural network
    #
    #Parameters:
    # hidden_unit_length_list - Array of hidden layer units
    # input_unit_count - integer, number of input units (features)
    # output_class_count - integer, number of output classes
    #
    #Returns:
    #
    #Array of theta arrays randomly initialized to from -.5 to .5
    #

```

```

if length( hidden_unit_length_list ) == 0
    hidden_unit_length_list = [2]
end

unit_count_list = [input_unit_count]
unit_count_list = [unit_count_list, hidden_unit_length_list]
unit_count_list = [unit_count_list, output_class_count]
layers = length(unit_count_list)

Theta_L = [ rand( unit_count_list[i], unit_count_list[i-1]+1 ) - .5 for i = 2:layers
end

function print_theta(Theta_L)
    # print_theta() is a helper function that prints Theta_L and architecture info
    # It does not actually "do" anything except print to stdout

    T = length(Theta_L)

    println()
    println( "NN ARCHITECTURE" )
    println( "$(T+1) Layers ($(T-1) Hidden)" )
    println( "$T Thetas" )
    println( "$(size(Theta_L[1],2)-1) Input Features" )
    println( "$(size(Theta_L[end], 1)) Output Classes" )
    println()

    println( "Units per layer (excl. bias unit)" )
    for t = 1:T
        if t == 1
            println( " - Input: $(size(Theta_L[t],2)-1) Units" )
        end
        if t < T
            println( " - Hidden $t: $(size(Theta_L[t],1)) Units" )
        else
            println( " - Output: $(size(Theta_L[t],1)) Units" )
        end
    end
    println()

    println( "Theta Shapes" )
    for l = 1:T
        println( "Theta $l: $(size(Theta_L[l]))" )
    end
    println()

    println( "Theta Values" )
    for t= 1:T
        println( "Theta $t:" )
        println ( Theta_L[t])
    end
    println()
end

function nn_cost(Y, Y_pred)
    #
    # nn_cost implements cost function for array inputs Y and Y_pred
    #
    # y is array of n_observations by n_classes
    # Y_pred is array with same dimensions as Y of predicted y values
    #
    if size(Y) != size(Y_pred)
        if size(Y,1) != size(Y_pred,1)
            error("Wrong number of predictions", "$(size(Y,1)) Actual Values. $(size(Y_pred,1)) Predicted Values")
        end
    end
end

```

```

        else
            error("Wrong number of prediction classes", "$(size(Y,2)) Actual Classes. $(size(Y,2)) Predicted Classes.")
        end
    end

    n_observations = size(Y,1)

    # Cost Function
    J = (-1.0 / n_observations) * sum((Y .* log(Y_pred)) + ((1-Y) .* log(1-Y_pred)))

    J

end

function nn_predict(X, Theta_L)
    #
    # nn_predict calculates activations for all layers given X and thetas in Theta_L
    # return all inputs and activations for all layers for use in backprop
    #
    # Parameters
    # X is matrix of input features dimensions n_observations by n_features
    # Theta_L is a 3D array where first element corresponds to the layer number, second to the layer weights, and third to the layer bias
    #
    # Returns
    # a_N - 1D Array of activation 2D arrays for each layer: Input (1), Hidden (2 to T), and Output (T+1)
    # a_Z - 1D Array of input 2D arrays to compute activations for all non-bias units
    # a_N[end] - 2D Array of predicted Y values with dimensions n_observations by n_classes
    #

    a_N = {}
    z_N = {}

    m = size(X,1)
    T = length(Theta_L)

    # Input Layer inputs
    push!(a_N, X) # List of activations including bias unit for non-output layers
    push!(z_N, zeros(1,1)) # add filler Z layer to align keys/indexes for a, z, and Theta

    # Loop through each Theta_List theta
    # t is index of Theta for calculating layer t+1 from layer t
    for t=1:T
        # Reshape 1D Array into 2D Array
        if ndims(a_N[t]) == 1
            a_N[t] = reshape(a_N[t], 1, size(a_N[t],1))
        end

        # Add bias unit
        a_N[t] = [ ones(size(a_N[t],1), 1) a_N[t] ]

        # Calculate and Append new z and a arrays to z_N and a_N lists
        push!(z_N, a_N[t] * Theta_L[t]') #'
        push!(a_N, sigmoid(z_N[t+1]))
    end

    z_N, a_N, a_N[end]

end

function back_prop(X_train, Y_train, Theta_L, lmda)
    #
    # Parameters

```

```

# X_train - Array of feature inputs with dimensions n_observations by n_features
# Y_train - Array of class outputs with dimensions n_observations by n_classes
# Theta_L is a 1D array of Theta values where 1D element is the layer number, the 2
# lmda - Float64 - lambda term for regularization
#
# Returns
# Y_pred as array of predicted Y values from nn_predict()
# Theta_Gradient_L as 1D array of 2D Theta Gradient arrays
#

n_observations = size(X_train,1)

T = length(Theta_L)

# Create Modified copy of the Theta_L for Regularization with Coefficient for bias u
# Create variable to accumulate error caused by each Theta_L term in layer a_N[n+1]
Theta_Gradient_L = {}
regTheta = {}
for i=1:T
    push!(regTheta, [zeros(size(Theta_L[i],1),1) Theta_L[i][:,2:]] )
    push!(Theta_Gradient_L, zeros(size(Theta_L[i])))
end

# Forward Pass
z_N, a_N, Y_pred = nn_predict(X_train, Theta_L)

# Backprop Error Accumulator
delta_N = {}
for t=1:T
    push!(delta_N, {})
end

# Error for Output layer is predicted value - Y training value
delta = Y_pred - Y_train
if ndims(delta) == 1
    delta = reshape(delta, 1, length(delta) )
end

# Loop backwards through Thetas to apply Error to prior Layer (except input layer)
# Finish at T-2 because start at 0, output layer is done outside, the loop and input

# Output Error
delta_N[T] = delta

# Hidden Layers Error
for t=0:T-2
    delta = (delta * Theta_L[T-t][:,2:]) .* sigmoidGradient(z_N[T-t])
    delta_N[T-t-1] = delta
end

# Calculate error gradients (no error in input layer)
# t is the Theta from layer t to layer t+1
for t=1:T
    Theta_Gradient_L[t] = Theta_Gradient_L[t] + delta_N[t]' * a_N[t] #'
end

# Average Error + regularization penalty
for t=1:T
    Theta_Gradient_L[t] = Theta_Gradient_L[t] * (1.0/n_observations) + (lmda * regTheta[t])
end

Y_pred, Theta_Gradient_L
end

```

```

function fit(X_train, Y_train, Theta_L, lmda, epochs)
#
#fit() calls the training back_prop function for the given number of cycles
#tracks error and error improvement rates
#
#Parameters:
# X_train - Array of training data with dimension n_observations by n_features
# Y_train - Array of training classes with dimension n_observations by n_classes
# Theta_L - 1D array of theta 2d arrays where each theta has dimensions n_units[lay
# epochs - integer of number of times to update Theta_L
#
#Returns
# Theta_L - 1D array of Theta arrays
# J_List - Array (length = epochs) of result of cost function for each iteration

J_list = zeros( epochs )

for i=1:epochs
# Back prop to get Y_pred and Theta gradient
Y_pred, Theta_grad = back_prop(X_train, Y_train, Theta_L, lmda)
# Record cost
J_list[i] = nn_cost(Y_train, Y_pred)
# Update Theta using Learning Rate * Theta Gradient
for t=1:length(Theta_L)
# Start with a large learning rate; need to update this to be more intelligent t
# Need to update to change learning rate based on progress of cost function
if i < 100
learning_rate = 5.0
else
learning_rate = 1.0
end
Theta_L[t] = Theta_L[t] - ( learning_rate * Theta_grad[t] )
end
end
#println("Cost $i: $(J_list[i])")
end

Theta_L, J_list

end

function XOR_test(hidden_unit_length_list, epochs)
#
#XOR_test is a simple test of the nn printing the predicted value to std out
#Trains on a sample XOR data set
#Predicts a single value
#Accepts an option parameter to set architecture of hidden layers
#

println( "Training Data: X & Y")

# Training Data
X_train = [1 1; 1 0; 0 1; 0 0] # Training Input Data
Y_train = [0 1; 1 0; 1 0; 0 1] # Training Classes
println( X_train )
println( Y_train )

# Hidden layer architecture
hidden_layer_architecture = hidden_unit_length_list

# Regularization Term
lmda = 1e-5

```



```

# Initialize Theta based on selected architecture
Theta_L = initialize_theta(size(X_train,2), size(Y_train,2), hidden_layer_architectu

# Fit
Theta_L, J_list = fit(X_train, Y_train, Theta_L, lmda, epochs)

# Print Architecture
print_theta(Theta_L)

# Print Cost
println("Cost Function Applied to Training Data: $(J_list[end])")

# Predict
X_new = [1 0;0 1;1 1;0 0]
println( "Given X:")
println("$X_new")
z_N, a_N, Y_pred = nn_predict(X_new, Theta_L)
println( "Predicted Y:")
println("$ (round(Y_pred,3)) ")

Y_pred

end
XOR_test (generic function with 1 method)

```

Out [8]: Y_pred = XOR_test([2], 5000)

In [14]: Training Data: X & Y

```

1      1
1      0
0      1
0      0

0      1
1      0
1      0
0      1

```

NN ARCHITECTURE
3 Layers (1 Hidden)
2 Thetas
2 Input Features
2 Output Classes

Units per layer (excl. bias unit)
- Input: 2 Units
- Hidden 1: 2 Units
- Output: 2 Units

Theta Shapes
Theta 1: (2,3)
Theta 2: (2,3)

Theta Values
Theta 1:
-4.244057658582526 8.011119225793868 -8.218268440922987
3.9453941972708795 8.156949365387408 -7.86130314531543

```
Theta 2:  
5.735531100620542      12.358587250630872      -11.947865946940489  
-5.7395079853749005    -12.366440373984625      11.955769589062232
```

```
Cost Function Applied to Training Data: 0.006301094016920786  
Given X:
```

```
1      0  
0      1  
1      1  
0      0
```

```
Predicted Y:
```

```
.997    .003  
.996    .004  
.003    .997  
.003    .997
```

```
4x2 Array{Float64,2}:
```

```
Out [14]: 0.997177    0.0028127  
          0.995939    0.00404581  
          0.00273109  0.997279  
          0.00298615  0.997025
```